

Sample Size, Relative Risk, and Interval Interpolation Using Evidence based Likelihood Methods

Sarah Fletcher
Vanderbilt University Biostatistics

December 16, 2011

Contents

1	Introduction	1
2	Evidential Sample Size	1
2.1	Examples of <code>ss.mean()</code>	2
3	Evidential Relative Risk	3
3.1	Examples of <code>rr_like()</code>	3
4	Interval Interpolation	4
4.1	Examples of interpolated intervals	4

1 Introduction

This package is based on the Likelihood framework and requires a basic knowledge of the evidential paradigm. For a brief introduction please refer to Royall [1] and Blume [2]. The Likelihood framework focuses on interpreting data as statistical evidence, and uses likelihood ratios to measure the strength of statistical evidence of one hypothesis over another.

2 Evidential Sample Size

The likelihood paradigm defines evidence as strong, weak, and misleading, and these probabilities are the foundations of evidential sample size estimation. Sample sizes are calculated such that the probability of strong evidence given N and k is less than the specified strong evidence bound. In other words $P(\text{StrongEvidence}|N, k) < 0.80, 0.85, 0.90$. These strong evidence bounds can be adjusted by the user but default to the above.

For the purpose of interpreting and communicating the strength of evidence, it is useful to divide the continuous scale of the likelihood ratio into descriptive categories, such as weak, moderate and strong evidence. Such a crude categorization allows a quick and easily understood characterization of the evidence for one hypothesis over another. Benchmark values of $k = 8$

and 32 have been suggested to distinguish between weak, moderate and strong evidence. The function `ss.mean()` defaults to values of $k=8$ and 32, but can also be specified by the user. The $P(WeakEvidence|N, k)$ can be interpreted as $1 - P(StrongEvidence|N, k)$, since it is assumed that the $P(MisleadingEvidence|N, k)$ is approximately 0. These sample size calculations are modeled after methods presented by Strug [3].

The function `ss.mean()` will return a data frame containing all combination of values of k , S (Strong evidence), and δ (Effect size), and the resulting sample size to detect strong evidence in favor of one hypothesis over another.

2.1 Examples of `ss.mean()`

```
> ss.mean(k=c(8,32),S=c(0.80,0.85), delta=c(.5,1.5))
```

	k	Strong.evidence	Effect.size	Sample.size
1	8	0.80	0.5	37.156254
2	8	0.85	0.5	44.196580
3	8	0.80	1.5	4.128473
4	8	0.85	1.5	4.910731
5	32	0.80	0.5	52.002551
6	32	0.85	0.5	59.779554
7	32	0.80	1.5	5.778061
8	32	0.85	1.5	6.642173

```
> ss.mean(k=c(8,20,32),S=c(0.80,0.85,0.90), delta=c(.2,.5))
```

	k	Strong.evidence	Effect.size	Sample.size
1	8	0.80	0.2	232.22659
2	8	0.85	0.2	276.22862
3	8	0.90	0.2	340.42677
4	8	0.80	0.5	37.15625
5	8	0.85	0.5	44.19658
6	8	0.90	0.5	54.46828
7	20	0.80	0.2	294.12536
8	20	0.85	0.2	341.24509
9	20	0.90	0.2	408.94786
10	20	0.80	0.5	47.06006
11	20	0.85	0.5	54.59921
12	20	0.90	0.5	65.43166
13	32	0.80	0.2	325.01594
14	32	0.85	0.2	373.62222
15	32	0.90	0.2	443.03191
16	32	0.80	0.5	52.00255
17	32	0.85	0.5	59.77955
18	32	0.90	0.5	70.88511

```
>
```

3 Evidential Relative Risk

For a more thorough understanding of the conditional likelihood model for relative risk see Blume (reference [2], p. 2589) and for the full derivation see Royall (reference [1], p. 165). In short, of interest is the $RR = (y1/n1)/(y2/n2)$, where $n1$ is the total number of i.i.d. *Bernoulli*(θ_1) observations required to obtain $n1 - y1$ failures and $n2$ is the total number of i.i.d. *Bernoulli*(θ_2) observations required to obtain $n2 - y2$ failures. Let $RR = \theta_1/\theta_2$ be the relative risk for a successful event in group 1 versus group 2.

The function `rr_like()` takes in values $y1$, $n1$, and $y2$, $n2$, and returns a likelihood object that can be plotted or used to find support intervals. The variables `lo` and `hi` specify the plotting region, and default to 0 and 6 respectively. The variable `points` is the number of points in $[0,1]$ to calculate likelihood (defaults to 1000), and the variable `scale`, scales the maximum likelihood to 1 (defaults to TRUE). If `plot(rr_like())` is called, a likelihood plot is returned with support intervals at $k=8$ and 32 , but can be specified otherwise.

3.1 Examples of `rr_like()`

```
> rr <- rr_like(15, 45, 5, 29,hi=10)
> interval(rr, 8)
```

```
$endpoints
[1] 0.8308308 5.7757758
```


```
$like
[1] 0.125
```

```
> interval(rr, 32)
```

```
$endpoints
[1] 0.6606607 8.5185185
```

```
$like
[1] 0.03125
```

```
> plot(rr)
>
```




4 Interval Interpolation


If the user specifies less than 1000 points, the 1/8 and 1/32 plotting intervals (or other specified intervals) will be interpolated from the likelihood object. The function `interpolate.likelihood()` does not need to be specified and the likelihood plotting function will call it if a user chooses to plot less than 1000 points. The intervals will then be plotted in the usual way.

4.1 Examples of interpolated intervals

```
> rr_10points<- rr_like(15, 45, 5, 29, hi=10, points=10)
> plot(rr_10points)
>
```



```
> rr_20points<- rr_like(15, 45, 5, 29, hi=10, points=20)
> plot(rr_20points)
>
```



References

- [1] Royall RM. *Statistical Evidence: A Likelihood Paradigm*. Chapman and Hall: London, 1997.
- [2] Blume JD, “Tutorial in Biostatistics Likelihood methods for measuring statistical evidence.”, *Statist. Med.* 21:2563–2599, 2002.
- [3] Strug LJ, Rohde C, Corey PN. “An Introduction to Evidential Sample Size Calculations” *The American Statistician* 61: 207–212, 2007.